


Stanford AIMI 2025 Summer Research Internship


ResNet Rebels



Problem Statement

Objective: Develop a reliable AI model to detect pneumonia from chest X-rays

- Train classifier on chest X-rays datasets
- Improve performance using label extraction tools
- Extend to bounding box localization

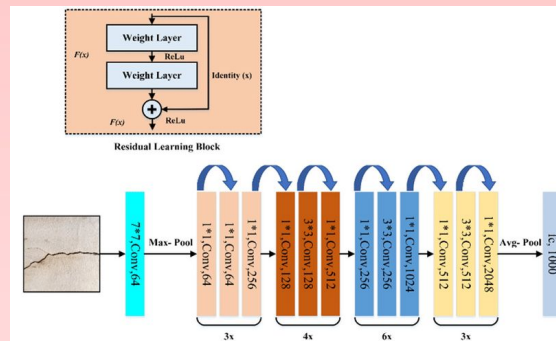


01 Classification Task



ResNet-18

- Shallow variant trained on ImageNet-1K
- Convolutional layers for feature extraction
- Pooling layers for dimensionality reduction
- Fully connected layers for classification across all classes

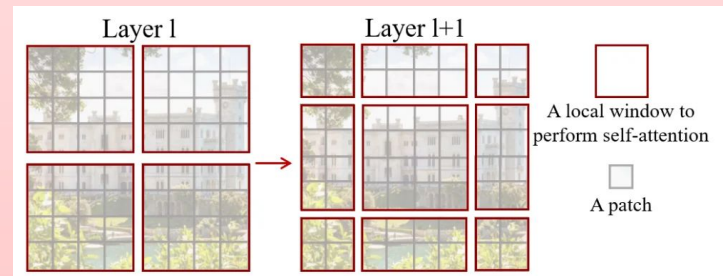
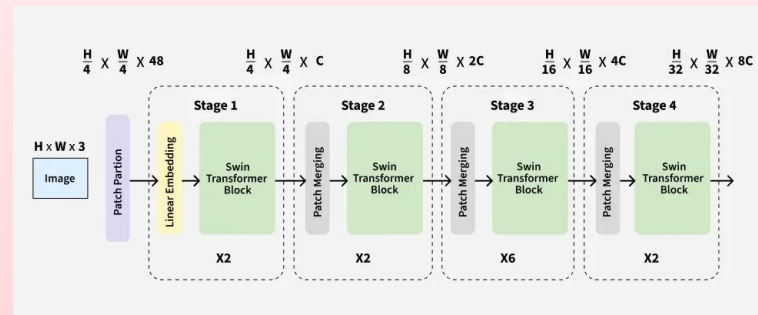


Epoch 1 | Loss: 0.4249 | Train Acc: 0.9517 | Train AUC: 0.9871 | Val Acc: 0.7962 | Val AUC: 0.8740
Epoch 2 | Loss: 0.1474 | Train Acc: 0.9965 | Train AUC: 1.0000 | Val Acc: 0.7642 | Val AUC: 0.8549
Epoch 3 | Loss: 0.0460 | Train Acc: 0.9983 | Train AUC: 1.0000 | Val Acc: 0.7832 | Val AUC: 0.8592
Epoch 4 | Loss: 0.0199 | Train Acc: 0.9996 | Train AUC: 1.0000 | Val Acc: 0.7729 | Val AUC: 0.8433
Epoch 5 | Loss: 0.0136 | Train Acc: 0.9987 | Train AUC: 1.0000 | Val Acc: 0.7953 | Val AUC: 0.8562

Evaluation metrics encouraged us to search for a better solution

Sliding Windows Transformer (Swin)

- Splits CXR images into fixed-size patches, reducing self-attention complexity from $O(N)$ to $O(N^2)$
- Allows cross-patch interaction across layers
- Four-stage architecture (96 → 192 → 384 → 768 channels) mimics CNN pyramids for downstream tasks
- Captures long-range context well, but demands more GPU memory


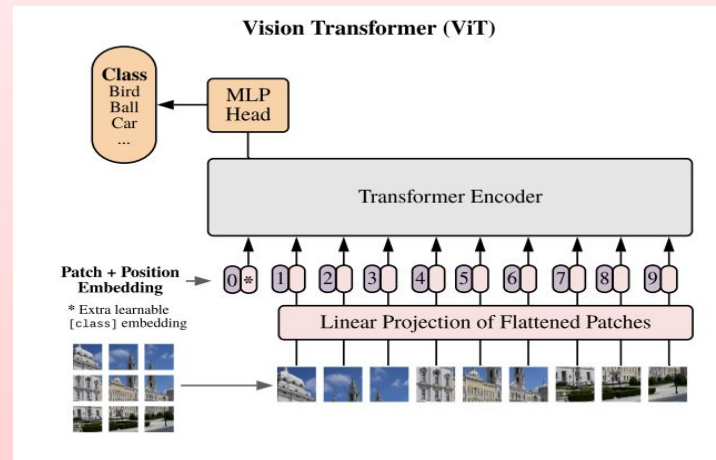


<https://arxiv.org/pdf/2103.14030>

<https://www.geeksforgeeks.org/computer-vision/swin-transformer>

Vision Transformer (ViT Base-16)

- Image → flattened 16×16 patches → linear layer → 196 patch tokens
- Add fixed position info + prepend [CLS] token → 197-token sequence
- **Transformer Encoder:** [CLS] attends to all patches via self-attention → builds global image summary
- **Classification Head:** Final [CLS] token passed to MLP → outputs class probabilities



<https://arxiv.org/abs/2010.11929>


<https://medium.com/analytics-vidhya/the-rise-of-attention-in-neural-networks-8c1d57a7b188>

Vision Transformer Pretraining

- ViT-Base pretrained on a corpus of ~0.5M CXRs
- Fine-tuned on the CheXpert dataset of 224,316 CXRs from the AIMI Center

Dataset Description

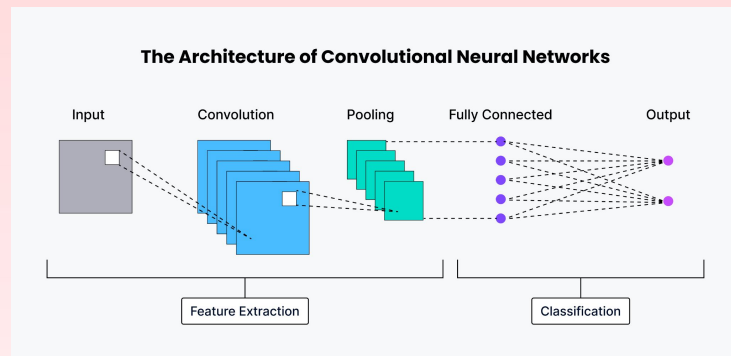
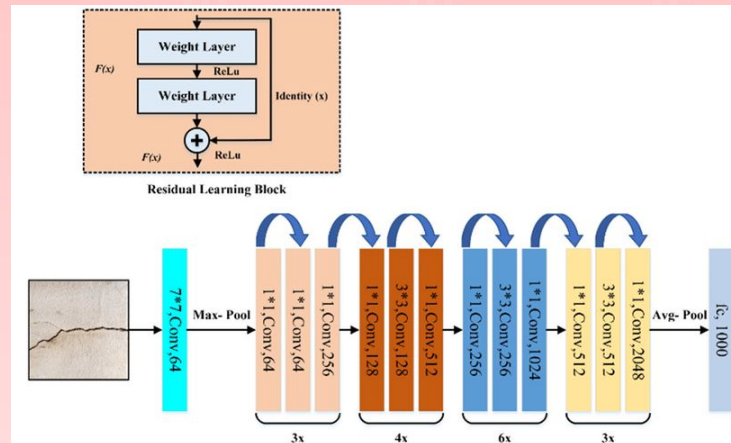
CheXpert is a dataset consisting of 224,316 chest radiographs of 65,240 patients who underwent a radiographic examination from Stanford Health Care between October 2002 and July 2017, in both inpatient and outpatient centers. Included are their associated radiology reports.



<https://aimi.stanford.edu/datasets/chexpert-chest-x-rays>
<https://stanfordmlgroup.github.io/competitions/chexpert/>

ResNet-18 with ImageNET Pretraining

- Shallow variant of the ResNet convolutional neural network trained on ImageNet-1K (2012).
- Convolutional layers extract features from an input image
- Pooling layers reduce the dimensions of data by the outputs of neuron clusters at one layer into a single neuron in the next layer.
- The fully connected layer connect every neuron in the penultimate layer to every neuron in the final layer (which correspond to all possible classes).
- ~11.7M parameters.

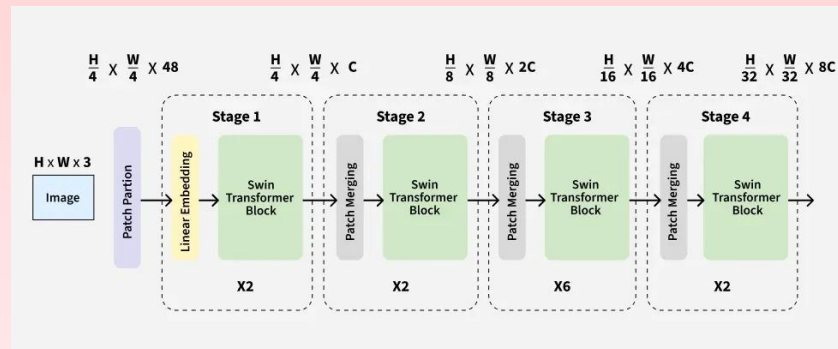
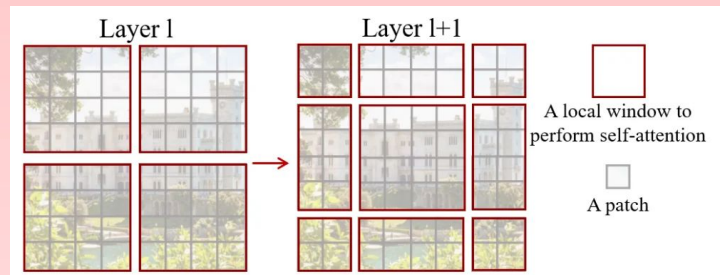


<https://roboflow.com/model/resnet-50>,

<https://zilliz.com/glossary/convolutional-neural-network>

Sliding Windows (Swin) Transformer

- Each layer cuts the CXR into smaller, fixed-size patches → lower computational cost ($O(N)$ vs. $O(N^2)$) of measuring pixel interactions via self-attention.
- The patches are slid in future layers so that information from multiple discrete windows can be leveraged in classification.
- Four hierarchical stages (sometimes $96 \rightarrow 192 \rightarrow 384 \rightarrow 768$ channels) mimic a CNN pyramid for easy use in detectors/segmenters.
- ~88 M parameters; needs more GPU memory but captures long-range context.

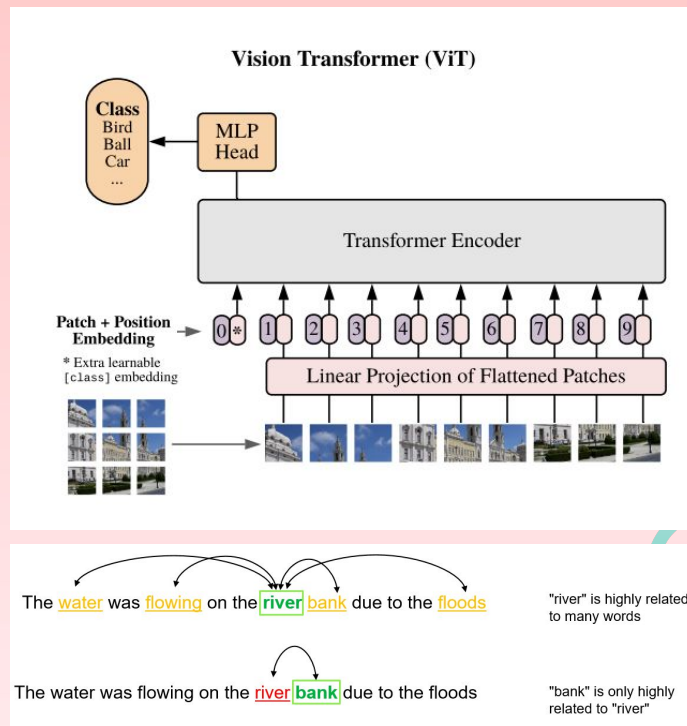


<https://arxiv.org/pdf/2103.14030>

<https://www.geeksforgeeks.org/computer-vision/swin-transformer/>

Vision Transformer (ViT-Base/16)

- Patch embedding: image \rightarrow flattened 16×16 patches \rightarrow linear layer \rightarrow 196 patch tokens.
- Positional + token addition: prepend the [CLS] vector, add fixed position encodings to every token \rightarrow 197-token sequence (box "0" in the figure is [CLS]).
- Transformer encoder: multi-head self-attention lets [CLS] attend to—and be attended by—every patch, accumulating a global summary of the input.
- The final [CLS] output feeds the small classifier MLP to produce class probabilities.
- Lowest parameter count (~8M).



<https://arxiv.org/abs/2010.11929>


<https://medium.com/analytics-vidhya/the-rise-of-attention-in-neural-networks-8c1d57a7b188>

Vision Transformer Pretraining

We leverage a ViT-Base pretrained on a corpus of ~0.5M CXRs and fine-tuned on the CheXpert dataset of 224,316 CXRs from the AIMI Center.

Dataset Description


CheXpert is a dataset consisting of 224,316 chest radiographs of 65,240 patients who underwent a radiographic examination from Stanford Health Care between October 2002 and July 2017, in both inpatient and outpatient centers. Included are their associated radiology reports.



<https://aimi.stanford.edu/datasets/chexpert-chest-x-rays>
<https://stanfordmlgroup.github.io/competitions/chexpert/>

Further Optimizations–Dropout

- Helps prevent memorizing the training data to improve generalization to unseen cases.
- Disables neurons randomly during training to encourage learning in many sub networks.
- In final configuration, standard dropout was applied with:
 - 10% probability in the main body (backbone) of the ViT.
 - A variable, tuned probability (between 25% and 50%) in the final classification layer (head).
- Stochastic depth dropout (where entire layers are randomly skipped) with 30% probability was applied.



<https://medium.com/data-science/dropout-in-neural-networks-47a162d621d9>

Dropout Influence

Validation performance
degrades ❌👎

Epoch 1	Loss: 0.4373	Train Acc: 0.8646	Train AUC: 0.9260	Val Acc: 0.8178	Val AUC: 0.8900
Epoch 2	Loss: 0.3491	Train Acc: 0.8899	Train AUC: 0.9530	Val Acc: 0.8092	Val AUC: 0.8784
Epoch 3	Loss: 0.3022	Train Acc: 0.9147	Train AUC: 0.9800	Val Acc: 0.7824	Val AUC: 0.8724
Epoch 4	Loss: 0.2475	Train Acc: 0.9478	Train AUC: 0.9822	Val Acc: 0.8057	Val AUC: 0.8749
Epoch 5	Loss: 0.1917	Train Acc: 0.9682	Train AUC: 0.9928	Val Acc: 0.8031	Val AUC: 0.8867

Without dropout

Reduced overfitting to training data ✅

Epoch 1	Loss: 0.6537	Train Acc: 0.7636	Train AUC: 0.8086	Val Acc: 0.7660	Val AUC: 0.8263
Epoch 2	Loss: 0.5972	Train Acc: 0.7775	Train AUC: 0.8475	Val Acc: 0.7686	Val AUC: 0.8539
Epoch 3	Loss: 0.5488	Train Acc: 0.7880	Train AUC: 0.8577	Val Acc: 0.7876	Val AUC: 0.8640
Epoch 4	Loss: 0.5241	Train Acc: 0.8141	Train AUC: 0.8880	Val Acc: 0.8031	Val AUC: 0.8817
Epoch 5	Loss: 0.4923	Train Acc: 0.8254	Train AUC: 0.8990	Val Acc: 0.8126	Val AUC: 0.8915


With dropout

Validation performance
improves ✅👍


Further Optimizations: Leveraging Additional Unlabeled Data

- **Challenge:** Our initial dataset was limited in size and unbalanced.
- **Solution:** We incorporated a second dataset of chest X-rays labeled using the CheXpert labeler on corresponding radiology reports.

Dataset 1




Dataset 2




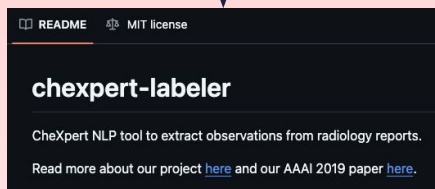
Task: convert
report to label.

Dataset 1 Label Distribution (N=2297)



CheXpert Labeling Setup

4233 Radiology
Notes 



4233 Labels 

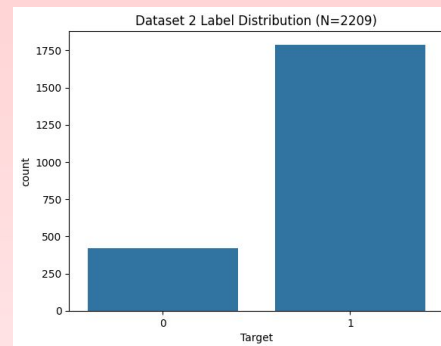
1 = pneumonia
0 = healthy
-1 = uncertain (less confident)

We tried:

- Marking uncertain cases as positive
- Marking uncertain cases as negative
- Removing uncertain cases

Removing uncertain cases yielded the largest, yet negligible (<0.1) AUROC improvement.

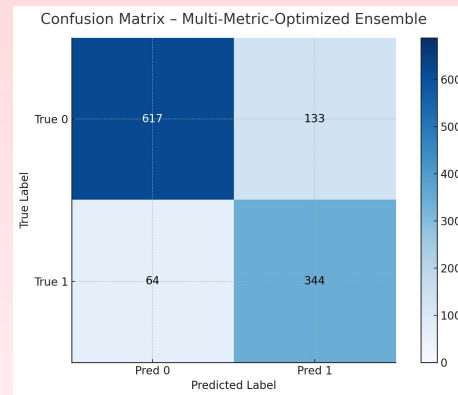
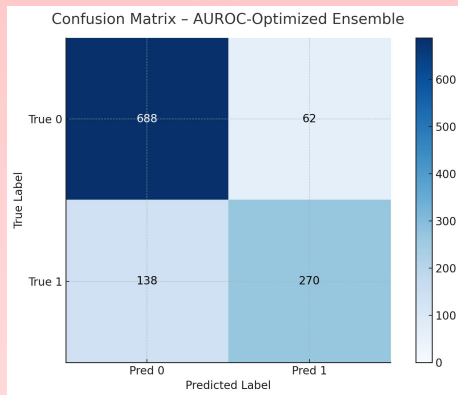
2209 additional (CXR, label)
samples used for training



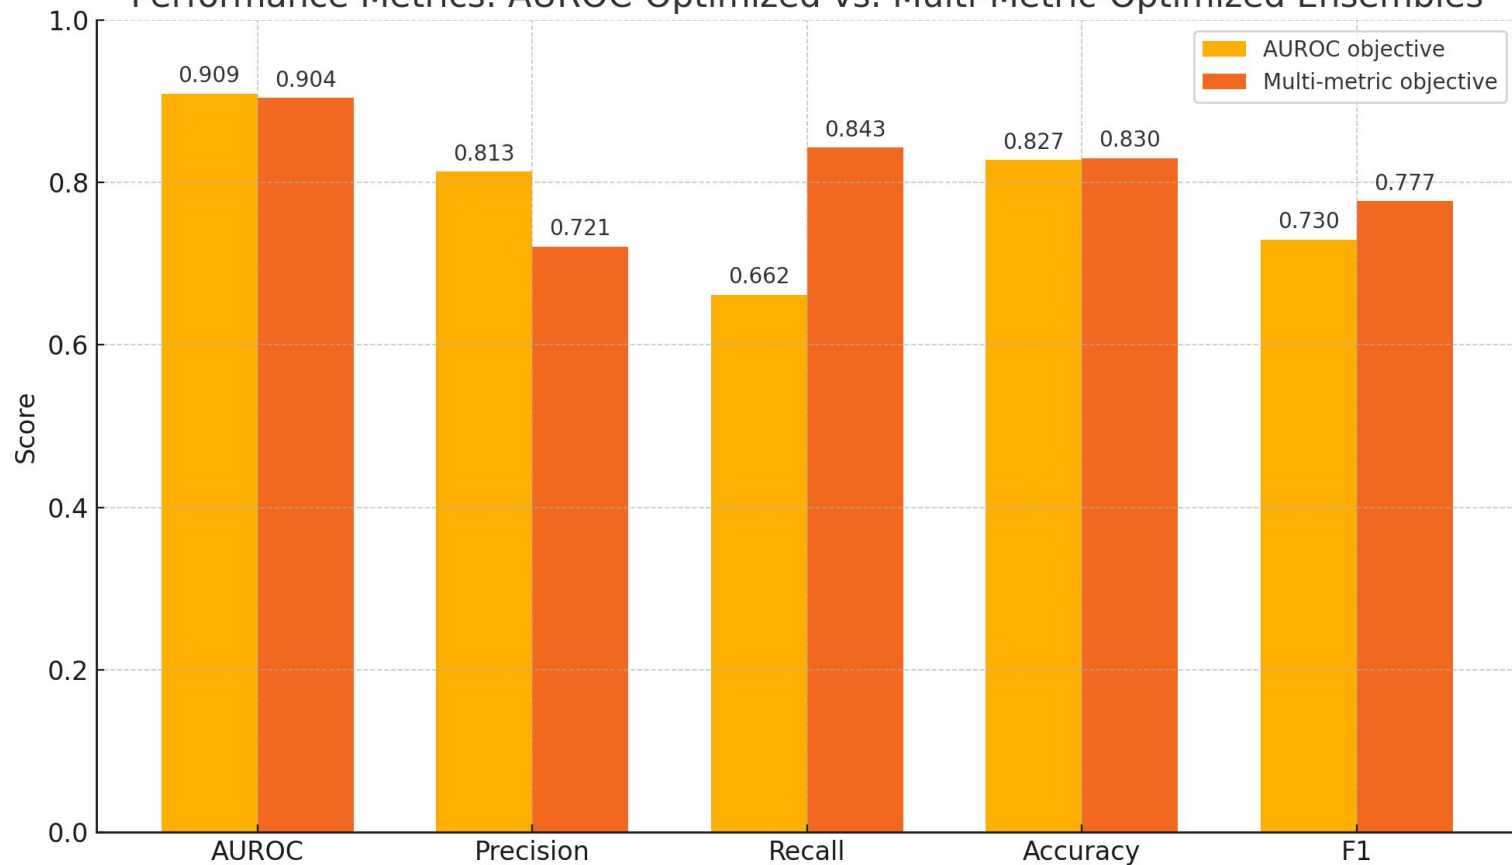
Further Optimizations: Hyperparameter Tuning

We used Optuna to efficiently test hundreds of combinations of hyperparameters.

- Initially tuned batch size, head dropout probability, learning rate, optimizer, and epochs with an objective of maximizing validation AUROC. We observed:
 - → low recall and F1 despite high AUROC.
 - → an alarming false negative count (see confusion matrices at right).
- Reran optimization with a weighted loss function that emphasized positive cases with a tuned weight; used an objective combining recall, f1, and AUROC.



Performance Metrics: AUROC-Optimized vs. Multi-Metric-Optimized Ensembles



Further Optimizations: Ensembling

- Keep 4 best configs from Optuna optimization.
- Train each 3× → 12 models
- Average logits
- Optuna picks best 6-model subset & threshold

Metric	Best Single ViT Tuned via. Optuna	6-ViT Ensemble	Δ
AUROC	0.898	0.904	+0.006
Recall	0.826	0.843	+0.017
F1	0.761	0.777	+0.016


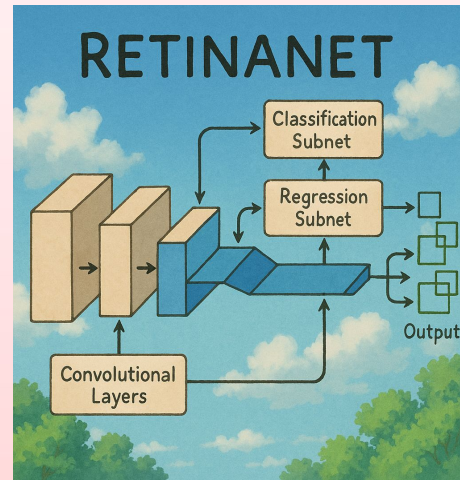
02

Localization Task



RetinaNet

- One stage detector with focal loss to address class imbalance
- ResNet50 backbone and FPN neck for multi-scale extraction
- Bounding box regression + classification
- Strong in small object detection


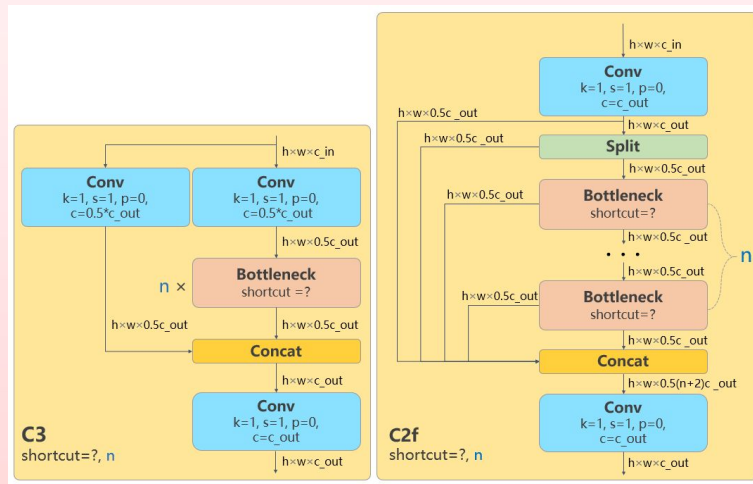


<https://paperswithcode.com/method/retinanet>

<https://blog.stackademic.com/training-a-retinanet-for-custom-object-detection-43fa6e518372>

YOLOv8/v11

- Tested model capacity by scaling from yolov8n → yolov8s → yolov8m
- Switched to YOLOv11
 - Higher mAP on COCO dataset with fewer parameters
 - ELAN-L creates more efficient gradient pathways
 - PGI provides clean path for learning signals to flow back



<https://docs.ultralytics.com/models/yolo11/>

https://mmyolo.readthedocs.io/en/latest/recommended_topics/algorithm_descriptions/yolov8_description.html

Performance Results

Best Run:

Metric	Value
Precision	0.75
Recall	0.69
F1 Score	0.72
mAP	0.47

Alterations:

- Trained on 1024×1024 images
- More epochs
- Heavy augmentations (mosaic, mixup, HSV)
- Optimized with AdamW + lr scheduling
- Early stopping & checkpointing

Final Pipeline

- Ran ViT classifier (pretrained on CheXpert) alongside YOLOv11 object detector
- If YOLO detects nothing but ViT is confident → fallback box with reduced confidence
- If YOLO detects pneumonia but ViT predicts negative → reduce YOLO box confidence
- Lowered confidence score when models disagreed
- Set output to “0” if both models predicted no pneumonia
- Merged final predictions for submission

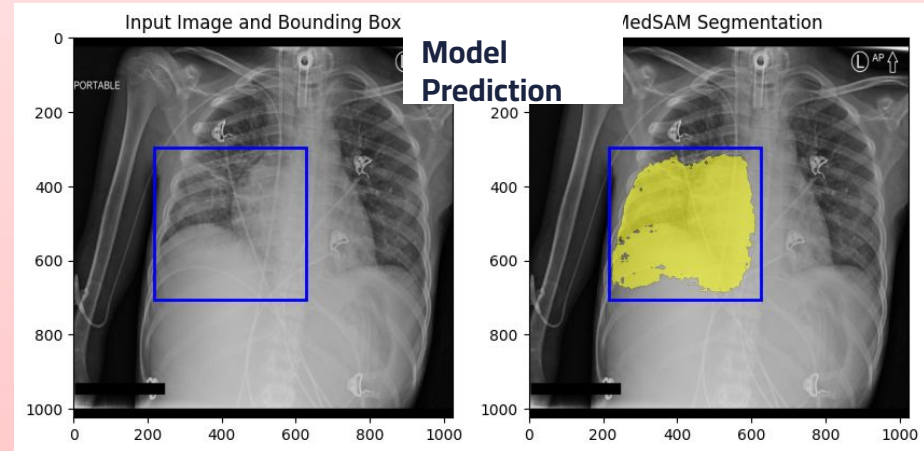

Results

	Metric	Value
0	Final Score	0.429844
1	Classification F1	0.644444
2	Classification AUC	0.753333
3	Localization Precision	0.020000
4	Localization Recall	0.009780
5	Mean IoU	0.539798
6	mAP@[.50:.95]	0.101250

- Combining lowered localization performance → need to explore other methods
- In the future:
 - Further hyperparameter tuning
 - Explore larger datasets
 - Implement further ensembling strategies
- Key limitations included small training set and limited hyperparameter tuning

Segmentation

- MedSAM - ViT Base
- Provide Model with Bounding Box of Ground Truth and Object Detection Prediction
- Transform Image and Convert to Tensors
- Turn Tensor into MedSAM Embedding



<https://colab.research.google.com/drive/19WNtRMbpsxeqimBlmJwtd1dzpalvK2FZ?usp=sharing>
<https://github.com/bowang-lab/MedSAM>

Insights & Reflections

- We realized that the lack of data that we had was a big issue
 - With more time + planning, we could have yielded better performance.
- We did have to switch versions of models
 - ex. Switching from YOLO v8 to YOLO v11
- For the final pipeline, we had to determine what to do if the classification model reached a different conclusion than the object detector.
- We were initially focused on the area under the curve, and realized that other evaluation metrics such as recall/sensitivity were extremely important in this setting as well.
- We tried many different approaches, and ended up taking different routes than we initially thought.



Thank you